

InnerLight Agency OS — Autonomous AI Agency on Your Desktop

One-line: Give it a complex brief; a 28-agent organization plans it, executes it on your machine, and proves every action was compliant.

THE PRODUCT OFFERING

A **Windows/Mac desktop application** (Electron) that runs a full autonomous AI agency locally. You type or speak a project brief — *"build a marketing site with a Supabase backend and a booking flow"* — and the system decomposes it, routes it through a hierarchical org of specialized agents, and **actually does the work**: writes code, runs commands, builds and tests software, browses the web, drafts content, and ships artifacts. Every action is intercepted, evaluated against policy, permitted or blocked, and recorded as a cryptographic proof of compliance (the OBEXGATE Witness Tetrad).

Three things set it apart from a chat assistant or from Manus.ai:

- 1. **It's an organization, not a single agent** — 28 specialized roles with enforced handoff discipline, QA gates, and a truth-verification layer.
- 2. **It runs on your machine** — Electron's main process gives native shell, filesystem, and process access. No rented cloud VM required for local work; your files, your tools, your control.
- 3. **Governance is in the execution path** — not a dashboard after the fact. Nothing risky happens without passing the enforcement gate, and every permitted action is witnessed and signed.

WHAT IT CAN DO (CAPABILITIES)

Capability	What that means in practice
Build software	Scaffold + write a website with backend, run the dev server, run tests, fix failures — full code lifecycle.
Operate the machine	Execute shell commands, manage files, run local tooling (git, package managers, build chains) with per-action approval.
Browse & research	Drive a real browser (Playwright), gather sources, verify claims, produce cited reports.
Spawn specialists	The orchestrator instantiates domain-specific agents on demand for whatever the brief requires.
Go-to-market	Draft content, marketing strategy, sales collateral, community responses — the Distribution layer.
Learn & advise	Analytics, experiments, customer-insight synthesis, strategy recommendations — the Learning layer.
Govern itself	Every action gated by OBEXGATE; loop limits, retry caps, and failure packets prevent runaway or silent failure.
Prove compliance	Each permitted action emits a signed Witness Tetrad (input + rule + decision + signature) — audit-ready.

THE COMPONENTS — FULL PARITY WITH THE 28-AGENT ORG

Every conceptual agent in the Dream Team library has a real, executable counterpart in the desktop app (`src/main/agents/`). "Full parity" = one running module per role, same hierarchy, same contracts.

Central Orchestrator	classifies the brief, routes, gates Learning output
└ Research Coordinator	→ Research Analyst · Market Intelligence · Idea Generator · Knowledge Librarian
└ Build Coordinator	→ Architect → UX Designer → Code Developer → QA Testing → Truth Agent
└ Operate Coordinator	→ DevOps · Data Pipeline · Security · Performance Optimization
└ Distribution Coord.	→ Marketing Strategy · Content Creation · Sales Enablement · Community Manager
└ Learning Coordinator	→ Analytics · Customer Insight · Experimentation · Strategy Advisor
└ Distribution Packager	assembles deliverable bundles

Three safety contracts run cross-cutting: failure-packet (no silent failures), trace-emitter (every handoff logged with monotonic sequence), loop-termination (hard cap of 15 iterations, bounded retries).

HOW IT'S BUILT (ARCHITECTURE)



- **Brain:** Anthropic Claude via `@anthropic-ai/sdk` (already a dependency).
- **Execution:** Electron main process — native computer access; optional E2B/cloud sandbox for untrusted or parallel work.
- **Backend:** Supabase (tickets, traces, artifacts, auth, RLS) with a local SQLite cache for speed.

- **Enforcement:** Orin runs as a sidecar service; the tool loop calls its gate at every action boundary. Reuse, don't rebuild.

PARITY SCORECARD VS. MANUS.AI

Dimension	Manus.ai	InnerLight Agency OS
Compute substrate	Rented cloud microVM	Your own machine (Electron) + optional cloud sandbox
Agent model	planner / executor (context isolation)	28-role org with enforced contracts
Tool loop	✔ ~50 calls/task	To build — main-process loop
Browser / shell / files	✔	Native via Node + Playwright — to wire
Autonomy trigger	✔	To build — scheduler/queue
Governance / proof	✗	✔ OBEXGATE enforcement + signed proof (the moat)

Status: brain ✔, org scaffold ✔, backend ✔, desktop shell ✔. To build: the tool-execution loop, native tool wiring, the Orin enforcement call, and an autonomous trigger. The hard, differentiating half — a governed, organization-shaped agent that proves its compliance — is what's already uniquely yours.